

Docker cookbook

Install

mac

```
1 brew install docker
2 brew install docker-compose
```

ubuntu

```
1 # remove the old docker
2 sudo apt-get remove docker docker-engine docker.io
3 # install repo
4 sudo apt-get update
5 # install packages to allow apt to use a repository over HTTPS:
6 sudo apt-get install apt-transport-https ca-certificates curl software-properties-common
7
8 # add Docker's official GPG key:
9 curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo apt-key add -
10
11 # add repo
12 sudo add-apt-repository "deb [arch=amd64] https://download.docker.com/linux/ubuntu
    $(lsb_release -cs) stable"
13
14 # install docker
15 sudo apt-get update
16 sudo apt-get install docker-ce
17
18
19 # install Compose 1.23.1
20 https://docs.docker.com/compose/install/#master-builds
21
22 sudo curl -L "https://github.com/docker/compose/releases/download/1.23.1/docker-compose-
    $(uname -s)-$(uname -m)" -o /usr/local/bin/docker-compose
23 sudo chmod +x /usr/local/bin/docker-compose
```

Try it

pull latest ubuntu

```
docker pull ubuntu:latest
```

start a container and login

```
docker container run -it ubuntu:latest /bin/bash
```

`docker container run` tells the Docker daemon to start a new container. The `-it` flags tell Docker to make the container interactive and to attach our current shell to the container's terminal. The image is latest ubuntu. The container will run `/bin/bash` when the container starts.

Docker will give your newly created container a random name if you don't assign one for it. Use `--name` to add name for it

```
docker container run --name webserver -it ubuntu:latest /bin/bash
```

Attaching to a running container

1. check which containers are running

```
docker container ls
```

2. attach a bash to it

```
docker container exec -it (name or id of the running container) bash
```

Dockerfile

Dockerfile is a plain-text document describing how to build an app into a Docker image.

```
1 FROM alpine
2 LABEL maintainer="xxxx"
3 RUN apk add --update nodejs nodejs-npm
4 COPY . /src
5 WORKDIR /src
6 RUN npm install
7 EXPOSE 8080
8 ENTRYPOINT ["node", "./app.js"]
```

build a image based on Dockerfile.

```
docker image build -t test:0.01 .
```

add a name and optionally a tag for the image. `.` means docker image builder should find a Dockerfile in the current location.

check the newly created image

```
docker image ls
```

| | | | | | |
|---|------------|------|--------------|---------------|--------|
| 1 | REPOSITORY | TAG | IMAGE ID | CREATED | SIZE |
| 2 | test | 0.01 | 551943a101d3 | 2 minutes ago | 68.1MB |

create a container on top of it

```
docker container run -d --name web1 --publish 8080:8080 test:0.01
```

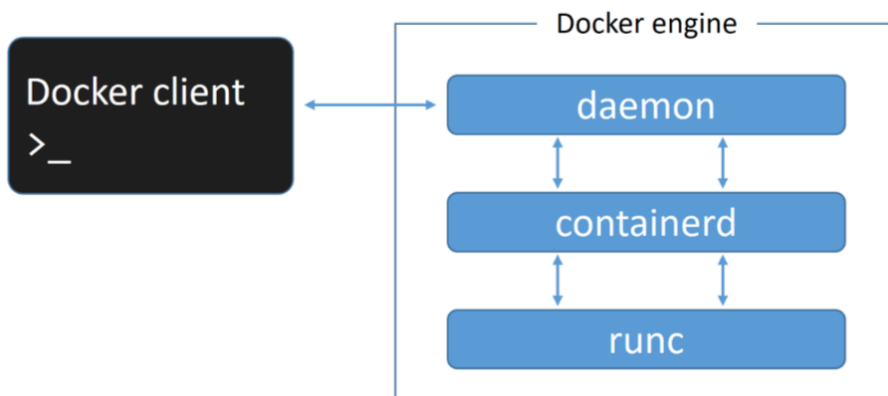
check if the new container is running

| | | | | | | | |
|---|--------------|-----------|-----------------|---------------|--------------|------------------------|-------|
| 1 | CONTAINER ID | IMAGE | COMMAND | CREATED | STATUS | PORTS | NAMES |
| 2 | 7cf67b047493 | test:0.01 | "node ./app.js" | 4 minutes ago | Up 4 minutes | 0.0.0.0:8080->8080/tcp | web1 |

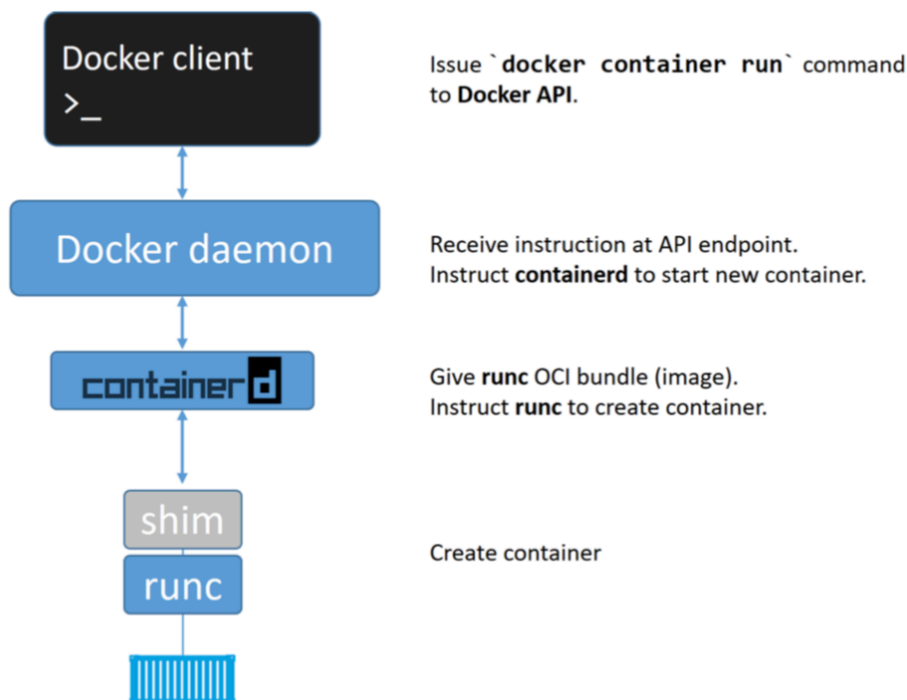
explanation of the arguments in the command:

1. `-d`: run container in background and print container ID
2. `--name`: assign a name to the container
3. `--publish` , `-p`: publish a container's port(s) to the host

docker engine



how they work together



docker image

Blueprint for the container.

- Image must contain all OS and application files required to run the app/service.
- Image don't contain a kernel — all containers running on a Docker host share access to the host's kernel.

- Docker image is just a bunch of loosely-connected read-only layers.

Cannot delete the image until the last container using it has been stopped and destroyed.

Docker images are stored in image registries.

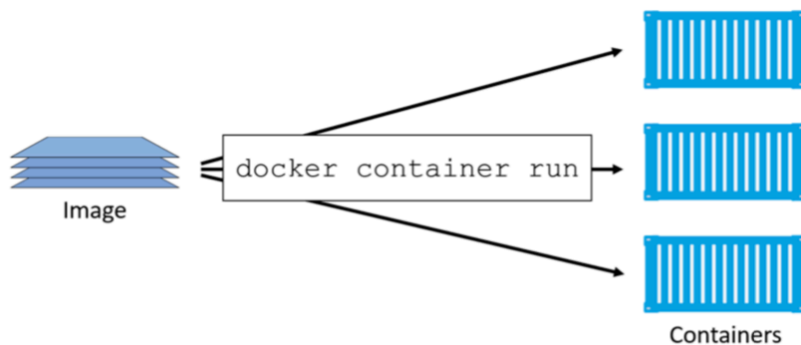
pull latest ubuntu

try `docker image pull ubuntu:latest`. If you do not specify an image tag after the repository name, Docker will assume you are referring to the image tagged as latest. However, `latest` is an arbitrary tag and is not guaranteed to point to the newest image in a repository. Add `-a` to pull all images in the repo.

- pull: `docker image pull <repository>:<tag>`
- remove dangling images: `docker image prune`
- filter: `docker image ls --filter key=value`
- format: `docker image ls --format "{{.Size}}"`
- search for repo: `docker search *****`
- inspect image: `docker image inspect`
- remove image: `docker image rm`
- remove all images: `docker image rm $(docker image ls -a -q) -f`
- build new image based on dockerfile: `docker image build -t test:0.01 .`
- push image to registry: `docker image push username/web:latest`

docker container

A container is the runtime instance of an image.



killing the main process in the container will also kill the container. Ctrl-PQ to exit the container without terminating it.

Example:

`docker container run -it ubuntu /bin/bash`

`-it` flags will make the container interactive and connect your current terminal window to the container's shell.

Containers run until the app they are executing exits.

Attaching to a container will create a new process.

lifecycle

- start a new container: `docker container run <options> <image>:<tag> <app>`
- start a stoped a container: `docker container start`
- stop a container: `docker container stop`
- remove a container: `docker container rm`
- reattach to a container: `docker container exec -it <container> bash`

restart policy

Enable Docker to automatically restart containers after certain events or failures have occurred.

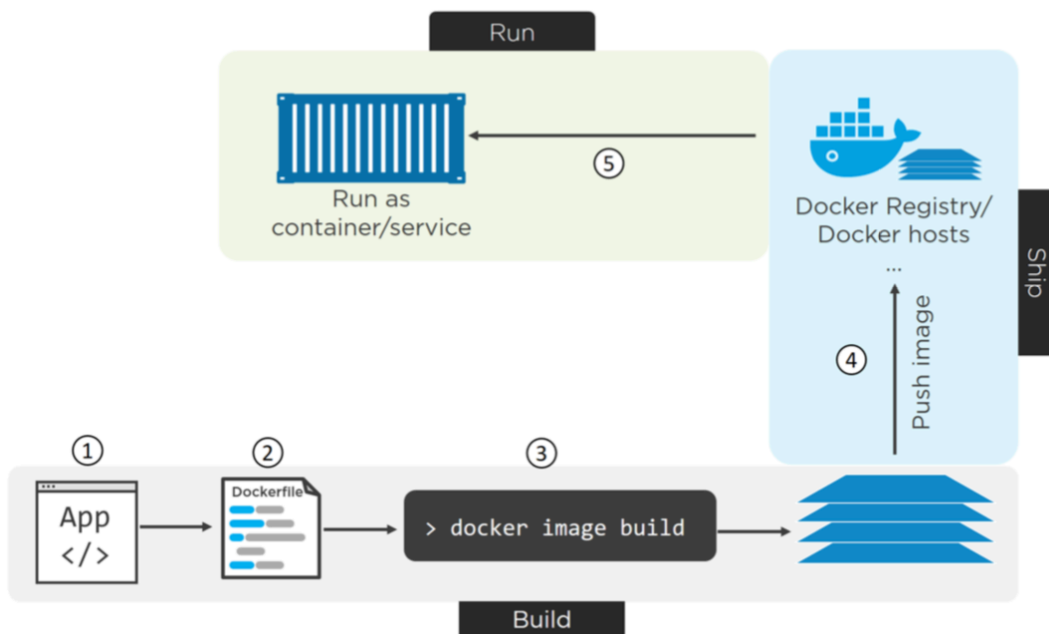
- `--restart always`: will always restart a stopped container unless it has been explicitly stopped, and the stopped container will be restarted when the Docker daemon starts.
- `--restart unless-stopped`: will always restart a container except when the container is in exited state.
- `--restart on-failure`: will restart a container if it exits with a non-zero exit code.

```
1 restart_policy:
2   condition: always | unless-stopped | on-failure
```

map port

`-p <host port>:<container port>`: maps ports on the Docker host to ports inside the container.

containerize an app



```
1 # base layer is alpine
2 FROM alpine
3 # add a label
4 LABEL maintainer="nigelpoulton@hotmail.com"
5 # RUN create a new temporary image as a layer with nodejs and npm installed in it
6 RUN apk add --update nodejs nodejs-npm
7 # copy files into a new image as a layer
```

```

8 COPY . /src
9 # set the working directory for the rest of the instructions in the file
10 WORKDIR /src
11 # install dependencies in the new image as a layer
12 RUN npm install
13 # export tcp 8080
14 EXPOSE 8080
15 # Entrypoint sets the command and parameters that will be executed first when a container is
  run.
16 ENTRYPOINT ["node", "./app.js"]

```

If an instruction is adding content such as files and programs to the image, it will create a new layer. If it is adding instructions on how to build the image and run the application, it will create metadata.

```
docker container run -d --name c1 -p 80:8080 web:latest
```

create a container named c1 that is based on web:latest image, and map 80 of host to 8080 of c1.

`docker image history <docker image>` checks how the image is built. Some instructions in the dockerfile create intermediate layers, sizes of the images are showed in the result.

General process to build an image

1. create a temporary container
2. run a line of the instructions in the dockerfile inside of the container
3. save the result as a new image layer
4. remove the temporary container

Optimization:

- combine multiple RUN commands as part of a single RUN instruction, glued together with double-ampersands (&&) and backslash () line-breaks

Multi-stage Builds

```

1 FROM node:latest AS storefront
2 WORKDIR /usr/src/atsea/app/react-app
3 COPY react-app .
4 RUN npm install
5 RUN npm run build
6
7 FROM maven:latest AS appserver
8 WORKDIR /usr/src/atsea
9 COPY pom.xml .
10 RUN mvn -B -f pom.xml -s /usr/share/maven/ref/settings-docker.xml dependency:resolve
11 COPY . .
12 RUN mvn -B -s /usr/share/maven/ref/settings-docker.xml package -DskipTests
13
14 FROM java:8-jdk-alpine
15 RUN adduser -Dh /home/gordon gordon
16 WORKDIR /static
17 # copy the build file to this image

```

```

18 COPY --from=storefront /usr/src/atsea/app/react-app/build/ .
19 WORKDIR /app
20 # copy the compiled file to this image
21 COPY --from=appserver /usr/src/atsea/target/AtSea-0.0.1-SNAPSHOT.jar .
22 ENTRYPOINT ["java", "-jar", "/app/AtSea-0.0.1-SNAPSHOT.jar"]
23 CMD ["--spring.profiles.active=postgres"]

```

The final image does not include `node:latest` and `maven:latest`, so it is much smaller.

The `ENTRYPOINT` specifies a command that will always be executed when the container starts.

The `CMD` specifies arguments that will be fed to the `ENTRYPOINT`. The `CMD` instruction in the dockerfile can be overridden by the the argument in the docker run command.

<https://docs.docker.com/v17.09/engine/userguide/eng-image/multistage-build/#before-multi-stage-builds>

cache

When docker excutes instructions in the dockerfile, it will try to find cache that matches the instruction. If the match can be found, docker will skip the instruction, link to that existing layer, and continues the build with the cache in tact. However, if any instruction results in a cache-miss, the cache is no longer used for the rest of the entire build.

Docker performs a checksum against each file being copied, and compares that to a checksum of the same file in the cached layer. If the checksums do not match, the cache is invalidated and a new layer is built.

`--no-cache=true` forces docker to ignore cache

Optimization:

- Try and build them in a way that places any instructions that are likely to change towards the end of the file.

squash

`docker image build --squash` can squashes all layers into one single layer. Squashed images do not share image layers

no-install-recommends

Use the `no-install-recommends` flag with the `apt-get install` command.

docker compose

Compose uses YAML files to define multi-service applications.

```

1 version: "3.5"
2 services:
3   web-fe:
4     build: .
5     command: python app.py

```

```

6     ports:
7       - target: 5000
8         published: 5000
9     networks:
10      - counter-net
11     volumes:
12      - type: volume
13        source: counter-vol
14        target: /code
15     redis:
16       image: "redis:alpine"
17       networks:
18         counter-net:
19
20 networks:
21   counter-net:
22
23 volumes:
24   counter-vol:

```

- version: defines the version of the Compose file format <https://docs.docker.com/compose/compose-file/compose-versioning/>
- services: defines the different application services
- networks: creates new networks
- volumes: creates new volumes

Each service will become a container.

`web-fe` will build the image that defined in the Dockerfile that is in the current location, then create a container and run `python app.py` as the main app. The port 5000(-target) in the container will be mapped to port 5000(published) inside the host.

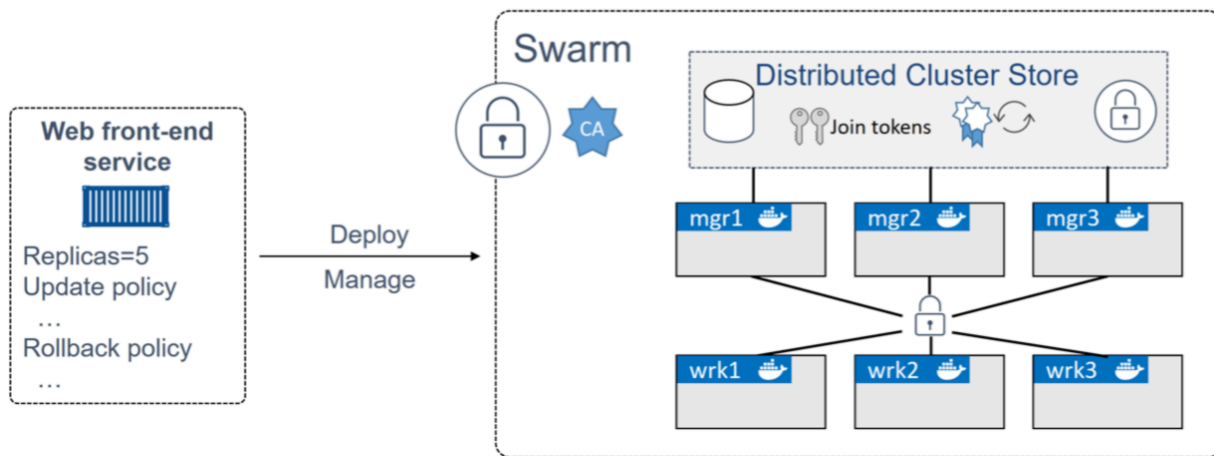
- `docker-compose up`: deploy a Compose app. and `&` to give the terminal back
- `docker-compose stop`: stop all of the containers in a Compose app without deleting them from the system.
- `docker-compose rm`: delete a stopped Compose app. It will delete containers and networks, but it will not delete volumes and images.
- `docker-compose restart`: restart a Compose app that has been stopped with `docker-compose stop`. If you have made changes to your Compose app since stopping it, these changes will not appear in the restarted app. You will need to re-deploy the app to get the changes.
- `docker-compose ps`: list each container in the Compose app.
- `docker-compose down`: stop and delete a running Compose app. It deletes containers and networks, but not volumes and images.

Swarm

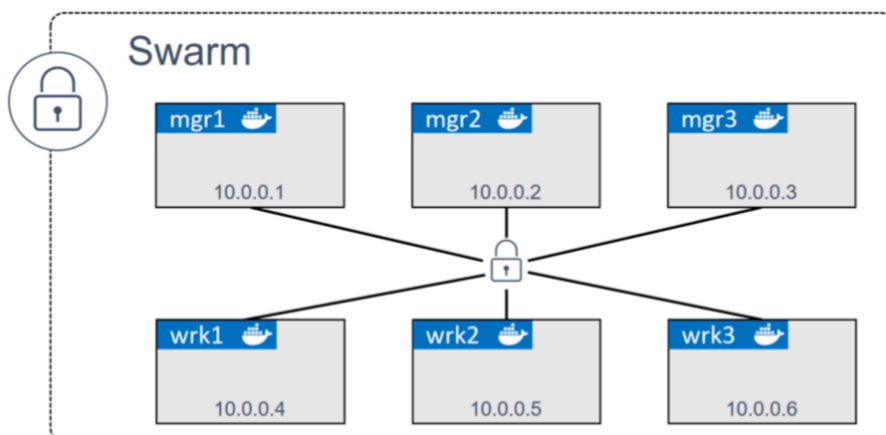
Docker Swarm is two things: an enterprise-grade secure cluster of Docker hosts, and an engine for orchestrating microservices apps.

a cluster has one or more Docker nodes. Nodes are configured as managers or workers. Managers look after the control plane of the cluster, meaning things like the state of the cluster and dispatching tasks to workers. Workers accept tasks from managers and execute them.

orchestration manages the lifecycles of containers.



build a secure swarm cluster



prerequisites:

- 2377/tcp: for secure client-to-swarm communication
- 7946/tcp and 7946/udp: for control plane gossip
- 4789/udp: for VXLAN-based overlay networks

1. create one manager and one worker

```
1 docker-machine create -d xhyve --xhyve-boot2docker-url  
https://github.com/boot2docker/boot2docker/releases/download/v18.06.1-ce/boot2docker.iso  
mgr1  
2 docker-machine create -d xhyve --xhyve-boot2docker-url  
https://github.com/boot2docker/boot2docker/releases/download/v18.06.1-ce/boot2docker.iso  
wrk1
```

2. log on the mgr1 and initialize swarm.

```
1 docker-machine ssh mgr1
2 #logged
3 docker swarm init \
4 --advertise-addr <ip of mgr1>:2377 \
5 --listen-addr <ip of mgr1>:2377
```

3. log on wrk1 and join the swarm as worker

```
1 docker-machine ssh wrk1
2 #logged
3 docker swarm join --token <token> <ip of mgr1>:2377 \
4 --advertise-addr <ip of wrk1>:2377 \
5 --listen-addr <ip of wrk1>:2377
```

4. check the swarm in the mgr1

```
1 docker node ls
```

The swarm should have multiple managers, and only one of them is "active" leader. The leader is the only one which can issue commands against swarm. If non-active managers receive commands for the swarm, they will proxy those commands to the leader.

Note: deploy odd number of managers. 3 or 5

prevent the restarted managers joining the cluster

`docker swarm update --autolock=true` will lock the swarm and return a token for the restarted managers to rejoin the cluster.

`docker swarm unlock`

- `docker swarm init`: initialize a new swarm and make this node the first manager
- `docker swarm join-token worker/manager`: show commands to join the swarm

service

`docker service create --name web-fe -p 8080:8080 --replicas 5 <image>`

The default replication mode of a service is replicated. This will deploy a desired number of replicas and distribute them as evenly as possible across the cluster.

The other mode is global, which runs a single replica on every node in the swarm.

- `docker service ps <service name>`
- `docker service scale <service name>=10`: change the number of containers among nodes
- `docker service rm <service name>`: remove the service. containers among the nodes will be removed too
- `docker service logs <service name>`
- `docker node ls`

Networks

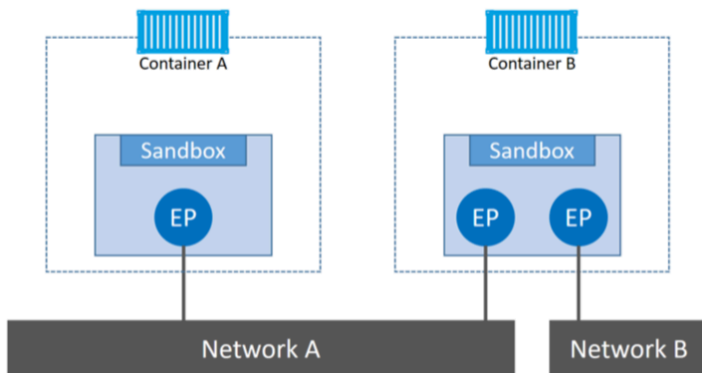
There are three default networks: bridge, null, host. and one virtual switch dockero

Container Network Model

design guide for Docker networking.

<https://github.com/docker/libnetwork/blob/master/docs/design.md>

- Sandboxes
Isolated network stack. It includes; Ethernet interfaces, ports, routing tables, and DNS config.
- Endpoints
virtual network interfaces that connect a sandbox to a network.
- Networks
software implementation of an 802.1d bridge



Libnetwork

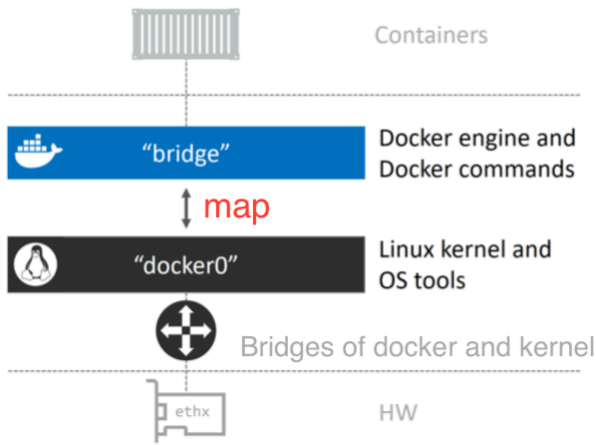
CNM implementation

Driver

Driver is in charge of the actual creation and management of all resources on the networks.

Single-host bridge networks

Every Docker host gets a default single-host bridge network. By default, this is the network that all new containers on this host will attach to unless other networks are specified.



create a network and attach a container to it:

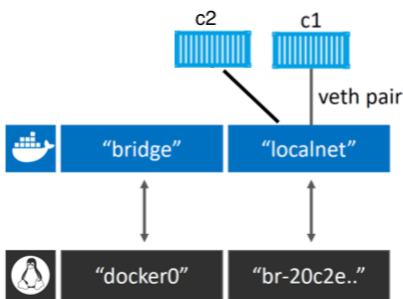
1. create a network

```
docker network create -d bridge localnet
```
2. create a container c1 which attaches to the network

```
docker container run -d --name c1 --network localnet alpine sleep 1d
```
3. create another container c2 which also attaches to the network

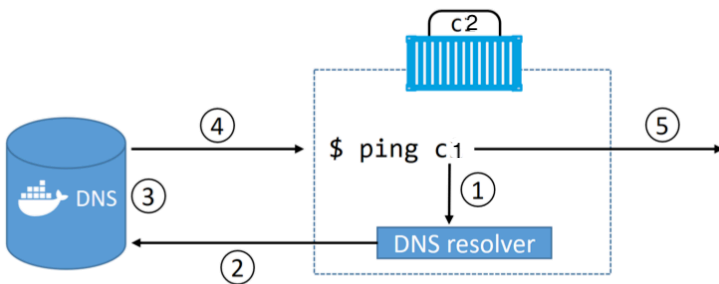
```
docker container run -it --name c2 --network localnet alpine sh
```
4. ping c1 from c2

```
ping c1
```



Name resolution:

c2 container is running a local DNS resolver that forwards requests to an internal Docker DNS server which maintains mappings for all containers. However, default bridge network on Linux does not support name resolution and name resolution only works for containers and Services on the same network..



1. The `ping c1` command invokes the local DNS resolver to resolve the name "c" to an IP address.

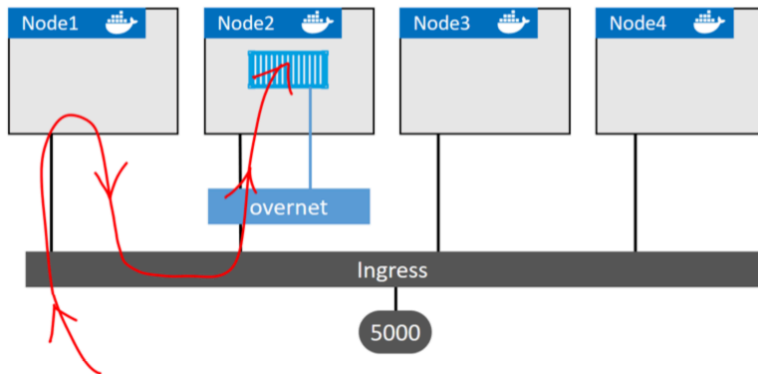
2. If the local resolver does not have an IP address for “c1” in its local cache, it initiates a recursive query to the Docker DNS server.
3. The Docker DNS server holds name-to-IP mappings for all containers created with the --name or --net-alias flags.
4. The DNS server returns the IP address of “c1” to the local resolver in “c2”.
5. The ping command is sent to the IP address of “c1”.

Port mappings let you map a container port to a port on the Docker host. Any traffic hitting the Docker host on the configured port will be directed to the container.

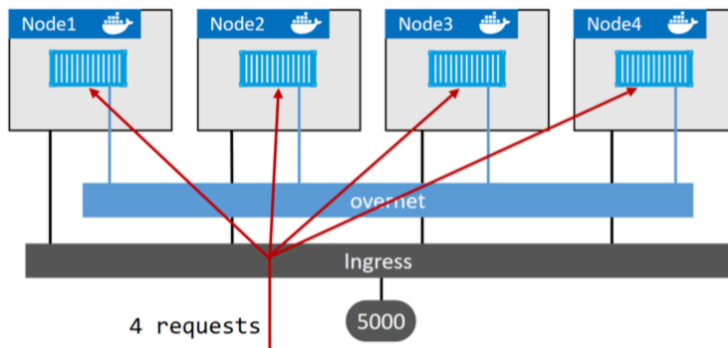
Load balancing

Services accessibility from outside of the cluster

- Ingress mode(default)
from any node in the Swarm



The traffic to 5000 of node 1 will be routed to node2 which the container is running in.



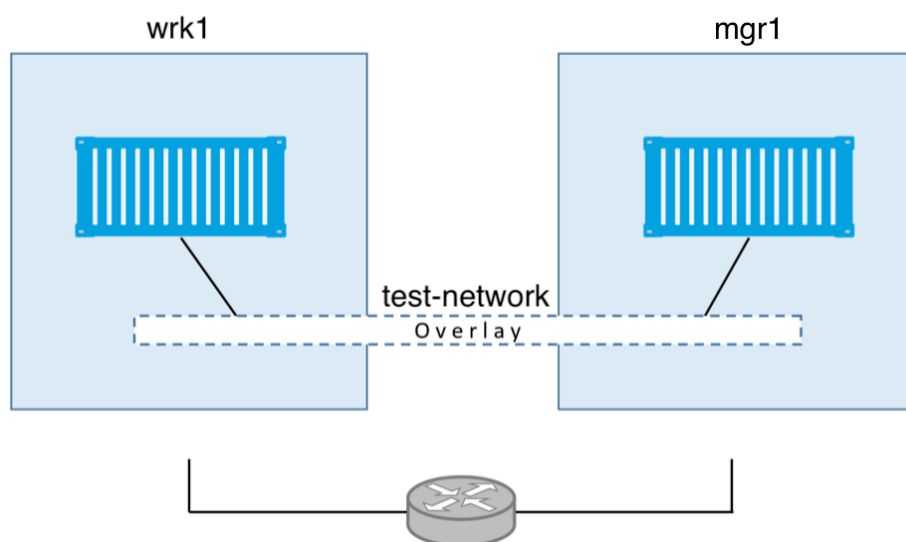
The traffic is balanced

- Host mode
nodes running service replicas
- `docker network create -d <driver> <name>`
- `docker network rm`
- `docker network inspect`

Overlay network

Build a overlay network for a swarm.

1. create a overlay network `test` on `mgr1`, `wrk1` cannot see it because it doesn't have containers that have attached to it yet.
`docker network create -d overlay test-network`
2. create a service that attaches to the new overlay network, now the `wrk1` can see the new overlay network
`docker service create --name test-service --network test-network --replicas 2 alpine sleep 1d`
3. get the ip of the newly created containers
`docker container inspect <container id>`
4. log on whichever container and ping another one
`docker container exec -it <container id> bash`
`ping xx.xxx.xx.xx`



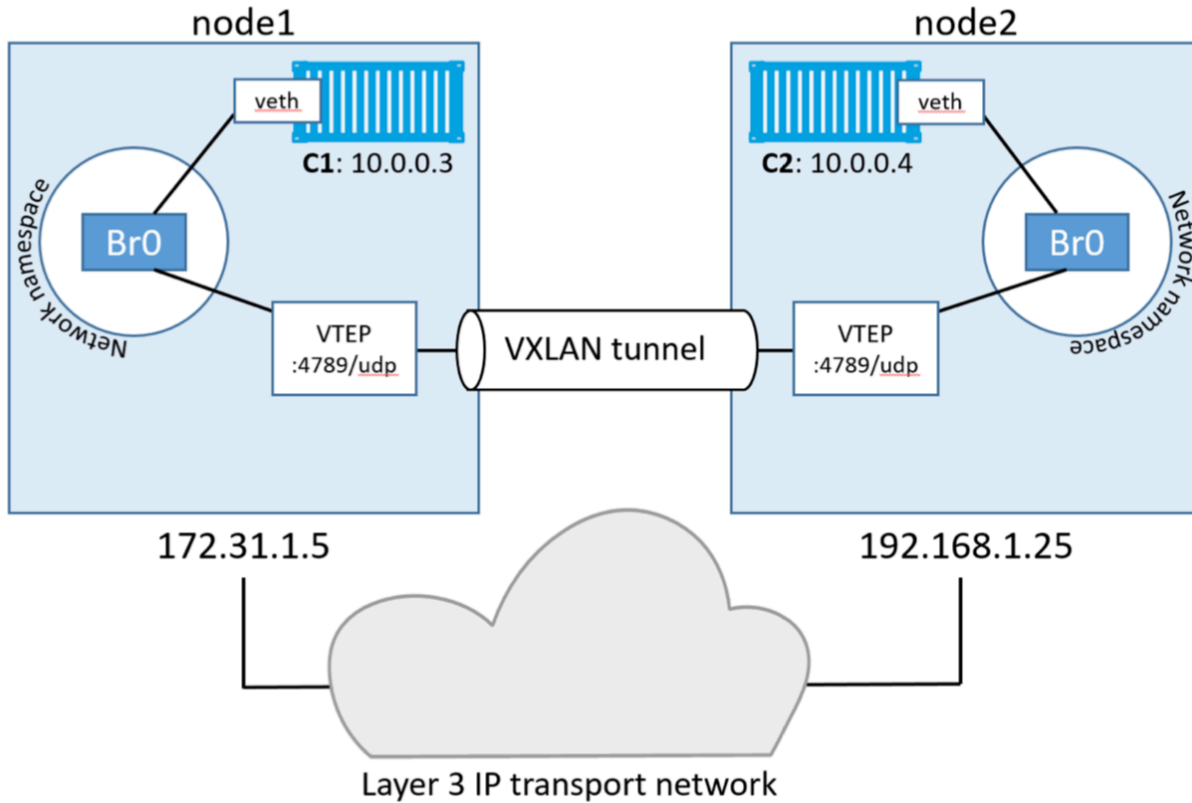
Docker overlay networking uses VXLAN tunnels to create virtual Layer 2 overlay networks. VXLANs let you create a virtual Layer 2 network on top of an existing Layer 3(IP) infrastructure.

Walkthrough:

1. When the VXLAN overlay network is created, docker creates a sandbox(a "container" for a network stack) on the host. A virtual switch (a.k.a. virtual bridge) called Bro is created inside the sandbox.
2. VTEP and veth?
3. c1 sends the first 'ping [xxx.xx.xx.xxx](#)' package to c2.
4. The traffic goes to Bro over veth. Bro doesn't know the mac address of c2, so it floods the packet to all ports.
5. The VTEP responds with the mac address of c2 back to Bro, so Bro update its mac address table.
6. Bro forwards all 'ping' packages to VTEP, and VTEP will encapsulate the Ethernet frame by adding VXLAN header into it.
7. VTEP on node1 sends the packages to VTEP on the node2. It decapsulates the packages and sends them to

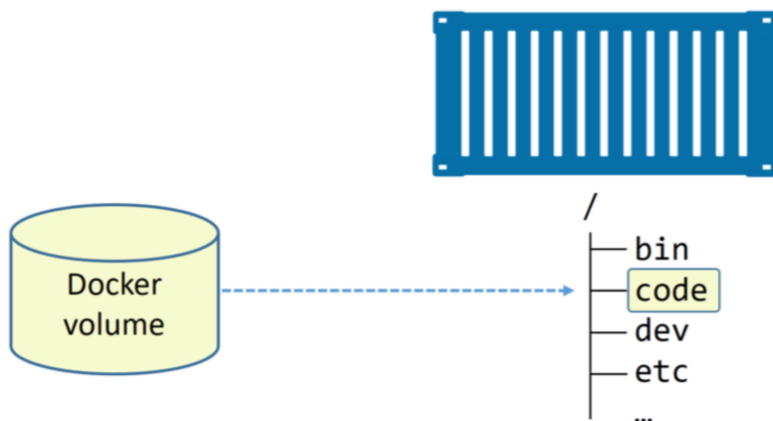
Bro on the node2.

8. Bro forwards packages to c2 to process.



Volumes

Docker creates a volume, then creates a container, and docker mounts the volume into it. The volume gets mounted to a directory in the container's filesystem, and anything written to that directory is written to the volume. If you then delete the container, the volume and its data will still exist. It's not possible to specify the host directory portion in a Dockerfile. This is because host directories are, by nature, host-dependent, meaning they can change between hosts and potentially break builds.



Different drivers

Mount a volume to a container

```
docker container run -dit --name voltainer --mount source=myvol,target=/vol alpine
```

Mount a volume to a service

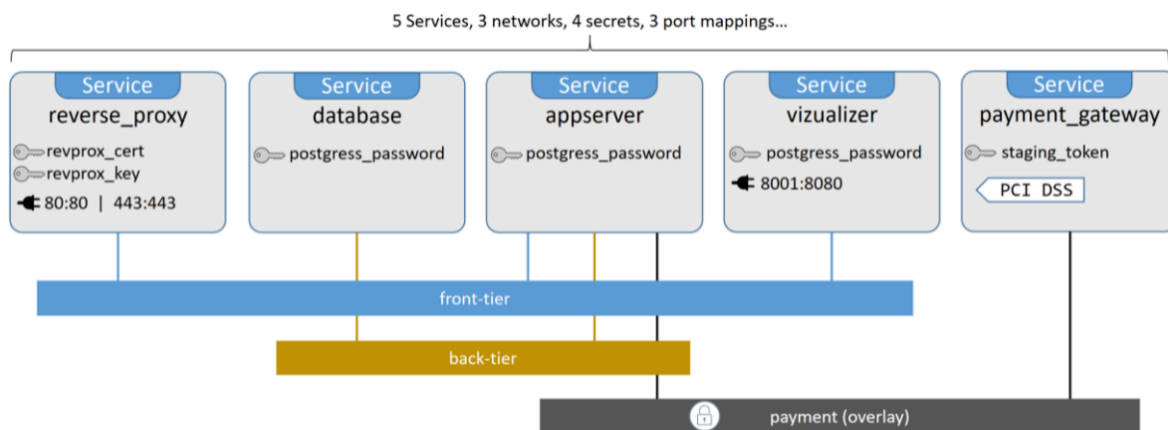
```
docker service create --name hellcat --mount source=myvol,target=/vol alpine sleep 1d
```

- `docker volume create myvol`: docker creates new volumes with the built-in local driver.
- `docker volume prune/rm`: remove the unmounted volumes or remove volumes by id

Docker stack

Docker stack simplified application management by providing; desired state, rolling updates, simple, scaling operations, health checks...

Good example: <https://github.com/dockersamples/atsea-sample-shop-app>



Generate secret

Docker can centrally manage sensitive data and securely transmit it to only those containers that need access to it.

1. create a docker secret file which contains a secret string

```
printf "This is a secret" | docker secret create my_secret_data -
```
2. create a service which has the access to the secret key

```
docker service create --name test --secret my_secret_data alpine:latest
```
3. check the content of the secret

```
docker container exec $(docker ps --filter name=test -q) cat /run/secrets/my_secret_data
```

Network

```
1 networks:
2   front-tier:
3   back-tier:
4   payment:
5     driver: overlay
```



```
6     driver_opts:
7         encrypted: 'yes'
```

By default, front-tier and back-tier will be created as overlay networks by the overlay driver. The payment network requires a encrypted data plane.

Secret

```
1 secrets:
2     postgres_password:
3         external: true
4     staging_token:
5         external: true
6     revprox_key:
7         external: true
8     revprox_cert:
9         external: true
```

secret must already exist before the stack can be deployed.

Service

```
1 reverse_proxy:
2     image: dockersamples/atseasampleshopapp_reverse_proxy
3     ports:
4         - "80:80"
5         - "443:443"
6     secrets:
7         - source: revprox_cert
8           target: revprox_cert
9         - source: revprox_key
10        target: revprox_key
11     networks:
12         - front-tier
```

Docker Stacks doesn't support builds. Secrets get mounted into service replicas as a regular file. The name of the file will be whatever you specify as the target value in the stack file, and the file will appear in the replica under /run/secrets on Linux.

Database

```
1 database:
2     image: dockersamples/atsea_db
3     environment:
4         POSTGRES_USER: gordonuser
5         POSTGRES_DB_PASSWORD_FILE: /run/secrets/postgres_password
6         POSTGRES_DB: atsea
7     networks:
8         - back-tier
9     secrets:
10        - postgres_password
11     deploy:
```

```

12     placement:
13     constraints:
14         - 'node.role == worker'

```

The `environment` key lets you inject environment variables into services replica. database container will only be deployed on worker node.

Appserver

```

1 appserver:
2   image: dockersamples/atsea_app
3   networks:
4     - front-tier
5     - back-tier
6     - payment
7   deploy:
8     replicas: 2
9     update_config:
10      parallelism: 2
11      failure_action: rollback
12     placement:
13     constraints:
14         - 'node.role == worker'
15     restart_policy:
16      condition: on-failure
17      delay: 5s
18      max_attempts: 3
19      window: 120s
20   secrets:
21     - postgres_password

```

`update_config` tells Docker how to act when rolling out updates to the service. Docker will update two replicas at a time and will perform a 'rollback' if it detects the update is failing. `restart-policy` tells Swarm how to restart replicas (containers) if and when they fail. Swarm will restart a replica if a replica stops with a non-zero exit code(on-failure). It will try to restart the failed replica 3 times, and wait up to 120 seconds to decide if the restart worked. It will wait 5 seconds between each of the three restart attempts.

Visualizer

```

1 visualizer:
2   image: dockersamples/visualizer:stable
3   ports:
4     - "8001:8080"
5   stop_grace_period: 1m30s
6   volumes:
7     - "/var/run/docker.sock:/var/run/docker.sock"
8   deploy:
9     update_config:
10      failure_action: rollback
11     placement:
12     constraints:

```

```
13     - 'node.role == manager'
```

`stop_grace_period` property overrides this 10 second grace period in which docker stops the container and waits to perform any clean-up operations. `volume` mounts pre-created volume or host dictionary into a service replic. `/var/run/docker.sock` is the IPC socket that the Docker daemon exposes all of its API endpoints on.

Payment_gateway

```
1 payment_gateway:
2   image: dockersamples/atseasampleshopapp_payment_gateway
3   secrets:
4     - source: staging_token
5       target: payment_token
6   networks:
7     - payment
8   deploy:
9     update_config:
10      failure_action: rollback
11     placement:
12       constraints:
13         - 'node.role == worker'
14         - 'node.labels.pcidss == yes'
```

`payment_gateway` replic will only be deployed on the worker node with `pcidss=yes` label.

deploy

- Create a new swarm with mgr1, wrk1, wrk2 and add a label to wrk2

....

```
docker node update --label-add pcidss=yes wrk2 add label into wrk2
```

- Create the secrets

```
openssl req -newkey rsa:4096 -nodes -sha256 -keyout domain.key -x509 -days 365 -out domain.crt
```

create docker secrets

```
docker secret create revprox_cert domain.crt
```

```
docker secret create revprox_key domain.key
```

```
docker secret create postgres_password domain.key
```

```
echo staging | docker secret create staging_token -
```

- deploy stack

```
docker stack deploy -c docker-stack.yml seastack
```

```
docker stack ps
```

```
docker stack rm
```

Log

check logs

```
docker logs <container>
```

`docker compose logs`